Text Mining

Abstract

(Lessions will be in English.) Lessons are divided between CS and DS + TTC students; the aim is to study and build search engines and recommendation systems. The first part is a presentations of Text Mining (an AI branch) with techniques of *Information Retrieval*, *Information Filtering*, *Text Classification* and *Summarization*, all using *open source* software. DS and TTC students will have labs using R and Python, CS students will use Java (Lucerne). The exam consists in a written test plus a project (done in groups of a maximum of 3 people), that can be followed by an oral test at will. On e-learning suggested books will be published.

Contents

Part I
Introduction and Text Representation

AI (Artificial Intelligence) was born in early '900: Machine Learning is only a small branch, started in '70s. Computational Linguistic and Statistics techniques are used as well. An Agent is considered *intelligent* if can understand the semantics of its Ambient, in that case texts written by humans in a *natural language*. To understand a text, the Agent must have a previous knowledge of the matter. Text Mining can be supported by Machine Learning but its techniques are more general and is built to emulate human behaviour. Other used techniques are DMs (Distributional Meanings) algorithms over time.

In 2004, Ian Witten (Weka project leader) published a paper titled "Text Mining" in which he reports that the first workshop was started in summer 1999: texts were processed in an automatic way to extract useful information.

Examples of Text Mining are *Sentiment* and *Opinion Analysis*; another example is *Text Summarization* (the ability to write a *snippet*, a short description of a text). *Reccomender Systems* are largely used by web-shopping platforms: they have revolutionized shopping and adverts suggesting useful information to users (positive filtering) or avoiding contents (negative filtering), analysing text content. Users are profiled by *Avatar*, their digital representations.

Text Analytics is used in healthcare to reduce the number of *fake news* on the web.

Knowledge Graphs are used to go behind traditional Text Mining techniques and have a better understanding of the text and improving the model.

The particularity of Text Mining is that a information in the text are not hidden but well written, and humans can reach it with no difficulty (but for text length); is largely used to analyze in a semi-automatic way lot of unstructured documents for decision making.

There are a lot of challenges in Text Mining:

- Text annotation;
- Dealing with large *corpora* and *streams*;
- Organizing semi- and un-structured data;
- Dealing with ambiguities on many levels (lexical, syntactic, semantic and pragmatic).

Information Retrieval aim is to find things on the Web, and has its roots in 70s. Techniques have to balance precision and efficiency.

Text Mining needs big data: good performance are reached only with a huge quantity of data due to algorithms complexity. For specific task contest-based, specific algorithms may have better performance or are less data-expensive.

Text analysis does not use databases but *corpora* (or *collections*): a group of unstructured documents; but sometimes information is well structured in the text (names or dates).

A *predictive analysis* of a text can recognise or detect a concept within a span of text, generally with supervised algorithms; examples are:

- opinion mining (positive or negative);
- sentiment analysis (fear, hope, desperation...);
- bias detection (pro/against);
- information extraction (is a person, is a place...);

- relation learning (is CEO of, is parent of...);
- text-driven forecasting (predicting events from twitter);
- temporal summarization (monitoring news or opinions about an event).

On the other hand, an *exploratory analysis* automatically discover patterns and trends of interest, generally done with an unsupervised approach.

1 Text and documents

To analyze a text, it have to be written so that a computer can easily read it: it have to be processed by a process that starts with *tokenization* (splitting text into units) and ends with the lexical analysis. Documents can be written in various ways, like free-format or in a semi-structued format (with descriptive and semantic metadata).

The first step consists in identifying input files and understand how to read them (creating a readable representation). The most used representation of text is the so called *bag of words*: introduced in 60s, each document is simply represented by a bag of words. Tokens are the basic features of the dataset, represented by the presence or absence (0 if absent, 1 if present, even if more times) in the document. The document is than processed to make it machine-readable. Another similar approach is to sign absolute frequency of words in the document. Both approaches do not consider words order, so "John is quicker than Mary" and "Mary is quicker than John" are represented in the same way.

Tokenization In this phase of the process, the document is divided in small units, each one representing a concept. Each unit (or *token*) is a sequence of characters with a semantic meaning (e.g. *"tree"*, *"San Francesco"* or *"01/01/1970"*). A token does not always correspond to a single word The task is difficult due to language-specific syntax and conventions, but *syntax-based* techniques have in general better performances. For example, in German word-splitting can boost performance in information retrieval tasks. Other oriental languages do not include spaces between words, so the tokenization process can return different solutions depending on the algorithm. Tokenization process is also context-depending (a social-media corpus contains a lot of abbreviations for example).
The process is made with a *parser* based on specific rules (defined with regular expressions) or statistical methods (machine learning or conditional probability).

Normalization It maps the text to query terms to the same form: sequence of characters with the same meaning have to be written in the same way (e.g. "USA $\rightarrow$ U.S.A."). The same concept is represented only in a single form: it is finalized to reduce tokens referred to the same concept. Uppercase and lowercase letters may be problematic due to semantic differences and language specific grammar. Thesaurus (or approximations like *WordNet*) that represent synonymy are largely used during the process. Spelling mistakes are managed using heuristics like phonetic pronunciation.

Stemming and Lemmatization In the stemming process words are "cutted" at the root which may be common with other words. In the lemmatization process instead words are transformed to their lemma. The second approach is more elegant but harder to implement (it

is extremely language-depending and it requires a thesaurus), in some cases benefits are not high enough to justify its use. Stemming on the other hand is faster and easier to implement, but it is based on the hypotheses that words with the same root have similar meanings, that is not always true (e.g. "author" and "authorize"); the task removes language specific suffixes, so algorithms are language depending. One of the most used stemming algorithm for English is Porter's algorithm, that uses simple morphological rules. Stemming and lemmatization do not always boost performance in English due to its syntax, but for other languages (Finnish) it works very well.

Stop words management   Stop words are very frequent words with no semantic meaning; they are generally removed from the text due to little importance in the analysis, but in some cases the text assumes a completely different meaning. Stop words are generally defined in a list (a *stop-list*, that can be taken from internet and modified to adapt it for the specific task) but may also include common and meaningless words. Search engines do include stop-words in the algorithm due to the risk of misunderstanding (but are managed in a special way).

Fundamental parts of the process are just tokenization and normalization: stemming, lemmatization and stop words management are context-depending. The algorithm is evaluated with two values: *precision* (the number of related documents depending on the query) and *recall* (the number of documents that match the query on the total of documents about the argument). The second one is fundamental in patents finding, the first one is important considering using-experience in web search engine (related documents must be presented as first results).

A different approach to the bag-of-words is the *N*-gram construction: *N* consecutive words are concatenated to catch linguistic expressions. This approach is more resources-expensive but it can boost performance a lot. Than a vocabulary is used to extract meaningful sequence of words (instances) or to exclude stop words from the removing process (the word can be a stop word if used alone, but in that context it assumes a particular meaning). Data is not stored on a matrix (due to the amount of memory required) but is stored using dynamic data structures (such as dictionaries and lists). *N*-grams tries to consider word order, storing more information than Bag of Words.

2   Zipf's Law

Originally used in Economics and than applied to Computational Linguistics, Zipf's Law (1949) describes the frequency of an event in a set. In particular, sorting words by frequencies in decreasing order, the product of use frequency and the rank order is approximately constant:

$$f(w) \propto \frac{1}{r(w)} = \frac{K}{r(w)}$$

where K is the corpus constant. So words can be divided in *head words* and *tail words* depending on the rank $r(w)$: head words are more frequent but in general are semantically meaningless.

According to this Luhn's Analysis, words do not describe document content with the same accuracy: word importance depends not only on frequency but also on position. So frequencies have to be weighted according to their position: two cut-offs are experimentally signed to discriminate very common or very rare words. Most significant words are those between the two cut-offs, those that are neither common nor rare. Weights are assigned according to two factors: word importance in the document and document importance in the collection.

$tf_{t,d}$ It represent the frequency of the term t in the document d, often normalized by document length:

$$w_{t,d} = \frac{tf_{t,d}}{|d|}$$

To prevent bias towards longer documents, it is not considered document length but the frequency of the highest frequency in the document:

$$w_{t,d} = \frac{tf_{t,d}}{\max\limits_{t_i \in d}\{tf_{t_i,d}\}}$$

idf_t It represent the inverse of the informativeness of the document for a term t:

$$idf_t = \log\left(\frac{N}{df_t}\right)$$

where df_t is the number of documents that contains t, an inverse measure of informativeness of t, and N is the total number of documents.

Weights are the product of the two indices:

$$w_{t,d} = tf.idf_{t,d} = \frac{tf_{t,d}}{\max\limits_{t_i \in d}\{tf_{t_i,d}\}} \cdot \log\left(\frac{N}{df_t}\right)$$

The weight w increases with the number of occurrences within a document and with the rarity of the term in the collection; this procedure is the best-known to calculate wights from 1983 (G. Salton et al.).

Part II

Natural Language Processing

To have a better representation of a text, N-grams are used to include a bit of semantics. Syntax is managed by two techniques: *Part of Speech Tagging* (POS) and *Entities Recognition*. The aim is to store identify and store positional information analyzing not N-grams but other data structures like pairs `<word, position>`.

Part of Speech Tagging   The aim of this procedure is to assign to each word its syntax role in the phrase. POS Tagging aims to make the Entities Recognition process easier, and can be managed in different ways. In addition, it can used to define rules in word generation (e.g. "a possessive adjective is followed by a name") and machine translation.

To have a perfect precision, a *lexicon* is required but it is not enough: it cannot manage ambiguity, that have to be solved following rules (also made by machine learning or hidden Markov models). This is the problem of this procedure: statistical and rule-based procedure have been proposed. Statistic methodologies are based on conditional probability (from a given *corpus*): probability of a category depending on the previous one is maximized to assign the output for the word. Rule based is in general more accurate but more time consuming because it often requires manual intervention. The process is built starting with short sentences to identify word categories and then longer sentences to identify rules.

Entities Recognition   The idea is to extract entities from text: it is a Information Extraction task. In this way the text representation is enriched with more information about its meaning. To complete the task, two approaches are identified: rule-based and statistics-based. According to the first approach, rules are used to identify if a (sequence of) word(s) can be an entity name; the second approach instead tries to maximize the probability with an hidden Markov model.

Message is sent by the sender and received by the recipient, who infers morphological information of the words, uses lexical rules to reconstruct the original message and compare it with its knowledge to understand its semantic (what is explicit, and does not change with the context) and pragmatic (what is implicit or context related). In the same way, a program have to rebuild the message and compare it whit a base of knowledge (usually modeled as a graph) to understand its content.

Anaphoras and cataphoras are constructions complex to manage correctly: to manage them, a *parser* is used to build a tree of the sentence (according on a formal representation of the grammar).

1 Topic Model

It is an unsupervised algorithm that clusterize documents according to their content. A topic is formal representation of the language used to write documents of a particular model, with a probability distribution for each token. In this way, it is easy to check which is the probability that a new document belongs to a topic. The same idea can be used using people own style

of writing instead of topic, and so on. This algorithm supports recommending systems that suggest contents to the user depending on how they write; it can also be used to generate new contents in automatic way or in machine translation.

2 Word Embeddings

For *word embedding* is indicated a set of NLP techniques to map a large namespace in a smaller one: in particular to transform concept space in a lower dimensional space in which it is possible to catch semantics similarity between words. Models to generate this mapping can be based on count of words or predictive models. This model can be used to check similarity between either documents or words, and to check topic similarity and words usage per topic.

To represent words as vectors, a $V \times V$ *1-hot* matrix is built in which each word has all values equal to 0 except for the corresponding column (that is equal to 1). This is a very sparse matrix, with no shared information between similar words. To reduce number of dimensions, not all words of the set V are used as columns: a small subset that includes just representative words is used as column. In this case, axis (columns) are putted manually and have a semantic meaning, but a previous knowledge of the domain is required. It is possible to catch relationships between words or to measure similarity between words using a distance index.

In a real scenario, it is not possible to catch all significant words in the domain (due to complexity of corpora), so automating algorithms are used. Results changes depending on the training corpus and its domain; often pre-trained (on general corpora such as dbpedia) models are retrained with domain-specific corpus to have a better representation of words. The model is trained trying to represent a word knowing its context (all other words in a defined *window* that set the size). So weights are taken counting how often a word co-occurs with another one. The number $f_{i.j}$ is the number of times that the word w_i appears in the context c_j: it is used to compute $P(w_i)$:

$$P(w_i) = \frac{to}{do}$$

The raw frequency is not a good representation: more used weights are frequency, tf-idf or PMI (*Pointwise Mutual Information*).

$$PMI(x, y) = \log_2 \left(\frac{P(x, y)}{P(x) \cdot P(y)} \right) \quad \in (-\infty; +\infty)$$

where x is the target word and y the context. To have better estimations, only positive estimation are taken:

$$PPMI(x, y) = \max(PMI(x, y), \ 0)$$

Rare words and *apax* produce higher scores, that risks to bias the analysis: probabilities of w and c are modified or a k-smoothing is used to reduce those scores.

$$P_\alpha(c) = \frac{count(c)^\alpha}{\sum_c count(c)^\alpha}$$

where (usually) $P_\alpha(c) > P(c)$ for rare c and $P_\alpha(c) < P(c)$ for common c. The second solution is the add-2 smoothing: a frequency of 2 is simply add to all elements of the matrix: effects are bigger with less frequent words, but high frequency words are little affected. From those high-dimensional matrices, dense lower-dimensional (usually around 300) vectors are learn as embeddings.

There are two mainly approaches to build an embedding: *count based* or *predictive models*. The first one uses statistic indices (based on frequencies like tf-idf) to map words in the new space; the second one tries to predict a word knowing its contest building an *auto-encoder* with a neural network.

2.1 Count based

This approach counts how many times two words (a word w and its neighbor c) co-occur in a large corpus; then the vector is mapped in a dense lower-dimensional space. Algorithms to reduce the number of dimensions are LDA (*Latent Dirichlet Allocation*), SVD (*Single Value Decomposition*) or GloVe (*Global Vectors for word representation*).

SVD is not really used due to its complexity (eigen values and vectors have to be computed for a very large matrix). From a mathematical point of view, the matrix X is decomposed in three matrices U (the matrix containing eigen vectors of XX'), S (a diagonal matrix containing eigen values of $X'X$) and V (the inverse of U) so that $X = USV'$. Eigen values σ_i (the element $S_{i.i}$) represent the variability of X captured by that eigen vector: removing them reduces the dimensions of the matrix. So, reducing dimensions of matrices to k (the selected eigen value that captures the desired percentage of variability) reduces dimensions of the matrix X minimizing information loss. In this way, words are mapped in a new space in which closer words are similar, but in which the original co-occurrence of words is lost (so it becomes a *black-box*). This algorithm is not really used due to its poor performance (decomposing matrices can be difficult, being $O(n^2)$ complex). Usually this algorithm is used only without stop-words and considering distance between words (using a *ramp window* to calculate proximity between words) or Pearson Correlation. Glove is used instead of SVD to solve its problems: it is trained only on non-zero elements in the matrix of the whole corpus. Its notation is similar to the previous one:

$$
\begin{aligned}
X \quad & \text{word-word co-occurrence matrix} \\
X_{i.j} \quad & \text{the frequency of } w_j \text{ occurring in the context } c_i \\
X_i = \sum_j X_{i.j} \quad & \text{frequency of a singular word} \\
P(j|i) = \frac{X_{i.j}}{X_i} \quad & \text{the probability of a word knowing its context}
\end{aligned}
$$

The ratio $\frac{P_{i.k}}{P_{j.k}}$ is larger then 1 if w_i is closer to w_k than w_j, lower otherwise and equal to 1 if the two words are equally distant from w_k. This ratio is used to modify the co-occurrence matrix: the quantity $\frac{P_{i.k}}{P_{j.k}}$ becomes a function $F(w_i, w_j, w_k)$ in which all variables w_i and w_j are word vectors and w_k the context. The problem is to find a valid F, but it can be simplified by

$$
F(w_i, w_j, w_k) = \frac{P_{i.k}}{P_{j.k}} = w_i' w_k + b_i + b_k = \log(X_{i.k})
$$

In this way however weights b are equally computed for both rare and frequent co-occurrences: the learning of vectors becomes a Least Squares Regression problem (with a weighting function):

$$J(\theta) = \sum_{i,k=1}^{V} f(X_{i.k})(w_i' w_k + b_i + b_k - \log(X_{i.k}))^2$$

$$J(\theta) = \frac{1}{2} \sum_{i,k=1}^{V} f(X_{i.k})(w_i' w_k - \log(X_{i.k}))^2$$

the weighting function f most common is:

$$f(x) = \begin{cases} \left(\frac{x}{x_{max}}\right)^{\alpha} & x < x_{max} \\ 1 & \text{otherwise} \end{cases}$$

but it can be any function that:

- $f(0) = 0$ and $\lim_{x \to 0} f(x) \log_2 x$ is finite;
- should be non-decreasing;
- should be small for big values of x.

This approach is fast to train, scalable and yields good performance.

2.2 Predictive models

Those models extract a subset of the corpus as train set, used to build an auto-encoder to map words in a new space with a Neural Network. The most popular algorithm belonging to this family is Word2Vec, introduced in 2013 to solve problems of SVD. It operate on a sample of the raw text, used to train a neural network that generate an auto-encoder of the text. Train data grows with the size of the window used to train the algorithm. The algorithm is used to predict the context knowing the original word (minimizing $J(\theta) = 1 - P(c|w)$ that maximize the probability of the prediction). Literature offers a better approach: the algorithm is used first to predict the context c (up to a certain level) giving a word w and then to predict the original word w knowing its context. Objective function (simplified) is the following:

$$J'(\theta) = -\prod_{t=1}^{T} \prod_{-m \leq j \leq m, j \neq 0} P(w_{t+j}|w_t)$$

$$= -\frac{1}{T} \sum_{t=1}^{T} T \sum_{-m \leq j \leq m, j \neq 0} \log P(w_{t+j}|w_t)$$

and then the prediction is made with a *naïve softmax*:

$$P(c|w) = \frac{\exp(u_c' v_w)}{\sum_{v=1}^{V} \exp(u_v' v_w)}$$

It captures more complex patterns than other approaches and can be used to other tasks; but on the other hand it does not scale well with corpus size and its very inefficient because it is trained only on a sample of the corpus. This algorithm is very context depending (the training is made on a corpus), and does not consider the changing of meaning over time.

Part III
Information Retrieval

Information Retrieval is an old problem in computer science: its main purpose is to find a documents on the Web. The main problem is the fact that documents are distributed between servers and not stored on the same computer; but also the quantity of data (and how to reach as many pages as possible) that are constantly updating can be very difficult to manage. Documents are managed in different ways, according to various criteria. Information Retrieval is done by *search engines* (that can search on the web - *web search engine* - or not).

Information Retrieval is a decision problem: how to identify and define the importance of information that satisfies the user. It is important to understand user's needs, to interpret the context and to yield documents that answer the query, according to its relevance. First search engines considered only topic relevance: it was not important to interpret a user query but just to yield relevance documents. Now the idea is to give a better user experience: the search process considers other variables such as geographical location and previous query done by the user; it can also be done without a user query (*Information Filtering Systems*) just filtering a stream of documents in real time.

In a search engine the query is usually not written in a formal language (but in a natural language): conditions have to be extrapolated from the text to reduce number of results. Information can be extrapolated from supports of different kinds: there are a lot of technical (due to different formats) and semantic (information is synthesized and represented preserving its meaning) problems.

A search engine is composed by two aspects (off-line and on-line): in the first part (offline) documents have to be indexed (and updated regularly by sub-parts of the archive) and represented in a formal way so that query can be done; than in the online part, the query is parsed to represent it in a logical language and than the match is made, according to classic set theory or using a similarity index on a vector space (cosine similarity, proximity and so on).